IST 782 Portfolio Milestone

M.S. Applied Data Science Syracuse University

Table of Contents

I. Introduction

Introduce Myself

Educational background

- B.B.A. in Operations and Supply Chain Management
- University of Wisconsin Whitewater, December 2017

Work experience

Position	Company	Location	Time
Solution Consultant	Smith & Nephew	Remote	Aug 2021 - Present
Global Demand Planner	Smith & Nephew	Remote	Nov 2020 – Aug 2021
Supply Chain Planner	Schreiber Foods	Green Bay, WI	Nov 2019 – Nov 2020
Global Demand Planner	Smith & Nephew	Austin, TX	Apr 2018 – Nov 2019
Supply Chain Intern	Mahindra USA	Houston, TX	May 2017 – Aug 2017

Professional accomplishments

- 4x dean’s list academic excellence award
- Co-treasurer and board member of campus APICS club
- 2nd place CSCMP supply chain case study competition
- APICS certified in production and inventory management (CPIM)
- Microsoft Excel Expert certification
- Amazon Warehouse Services (AWS) certified cloud practitioner
- Rapid Response certified contributor level 1
- Rapid Response certified author level 1

Skills

- R
- SQL
- Python
- Data Visualization
- Machine Learning
- Data Science
- Microsoft Excel
- Tableau
- Critical Thinking
- Team Leadership
- Public Speaking
- Analytic Problem Solving
- Alteryx
- Statistical Forecasting
- Supply Chain Analytics
- Amazon Web Services (AWS)
- Web Scraping

Why Data Science

The supply chain industry needs data scientists

- Statistical forecasting
- Inventory optimization
- Complex supply networks
- Manufacturing automation
- Risk management
- Real time data processing
- Predictive analytics

Data science is a growing field

- Per LinkedIn, there has been a 650% increase in data science jobs since 2012
- IBM says the demand for data scientists will continue to be strong for years
- The U.S. Bureau of Labor Statistics expects 11.5 Mil new data science jobs through 2026
- In 2020, data scientist was listed as the third best job in America according to Glassdoor

Flextronics Pulse Center, Silicon Valley

Statistics cited from https://towardsdatascience.com/is-data-science-still-a-rising-career-in-2021-722281f7074c

II. Project 1

Health and fitness database

Project Introduction

Goal: create a database for storing fitness related data such as exercise history and nutrition logs and use it to enhance the client experience

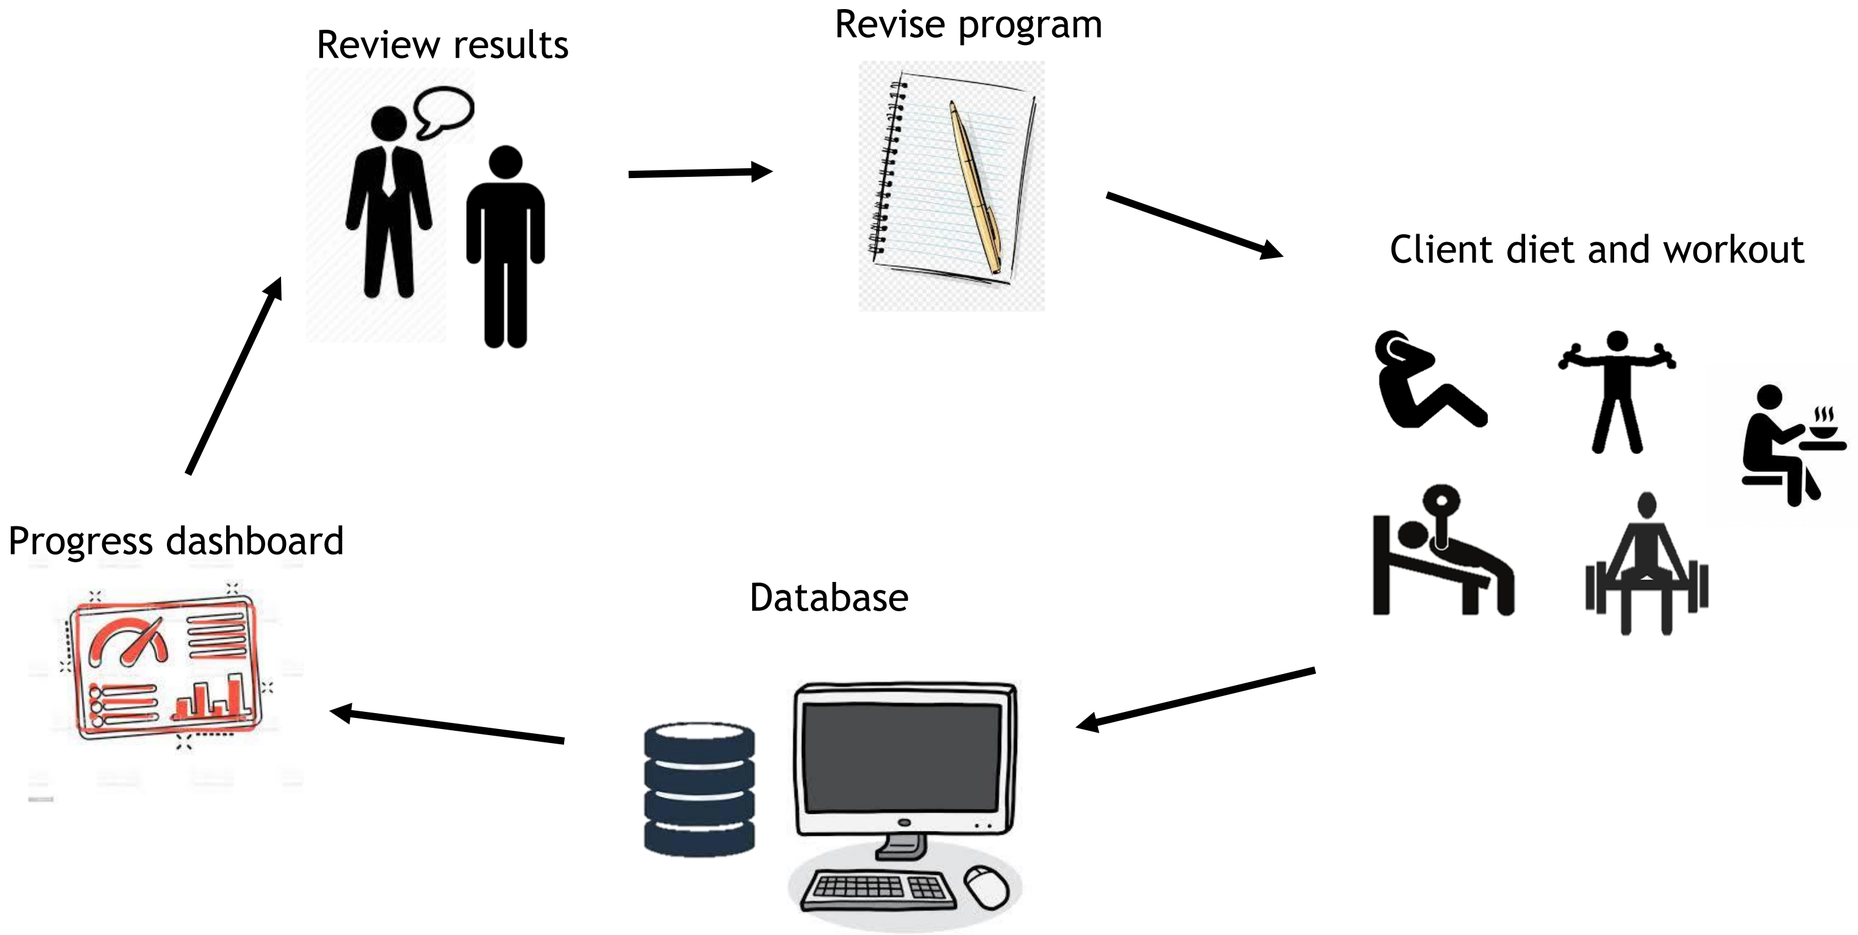

Database Development

CLIENT:A personal training client who is taking part in a weightlifting program and is using or did use the database for keeping track of their progression.

MEASUREMENT: a measurement taken by a client at a particular point in time

MUSCLE GROUP: a body part that is affected by lifting and can be measured

LIFT: a movement performed in a gym during a workout, that typically includes resistance, and impacts one or more muscle groups.

SET: a collection of repetitions

WORKOUT: a collection of sets

GYM: a facility that has the necessary equipment that can be used by a client to perform lifts

NUTRITION LOG: a diary that contains records of food or beverage items that are consumed

FOOD ITEM: an article of food or beverage that has calories and that is entered into the nutrition log

Entity Relation Diagram

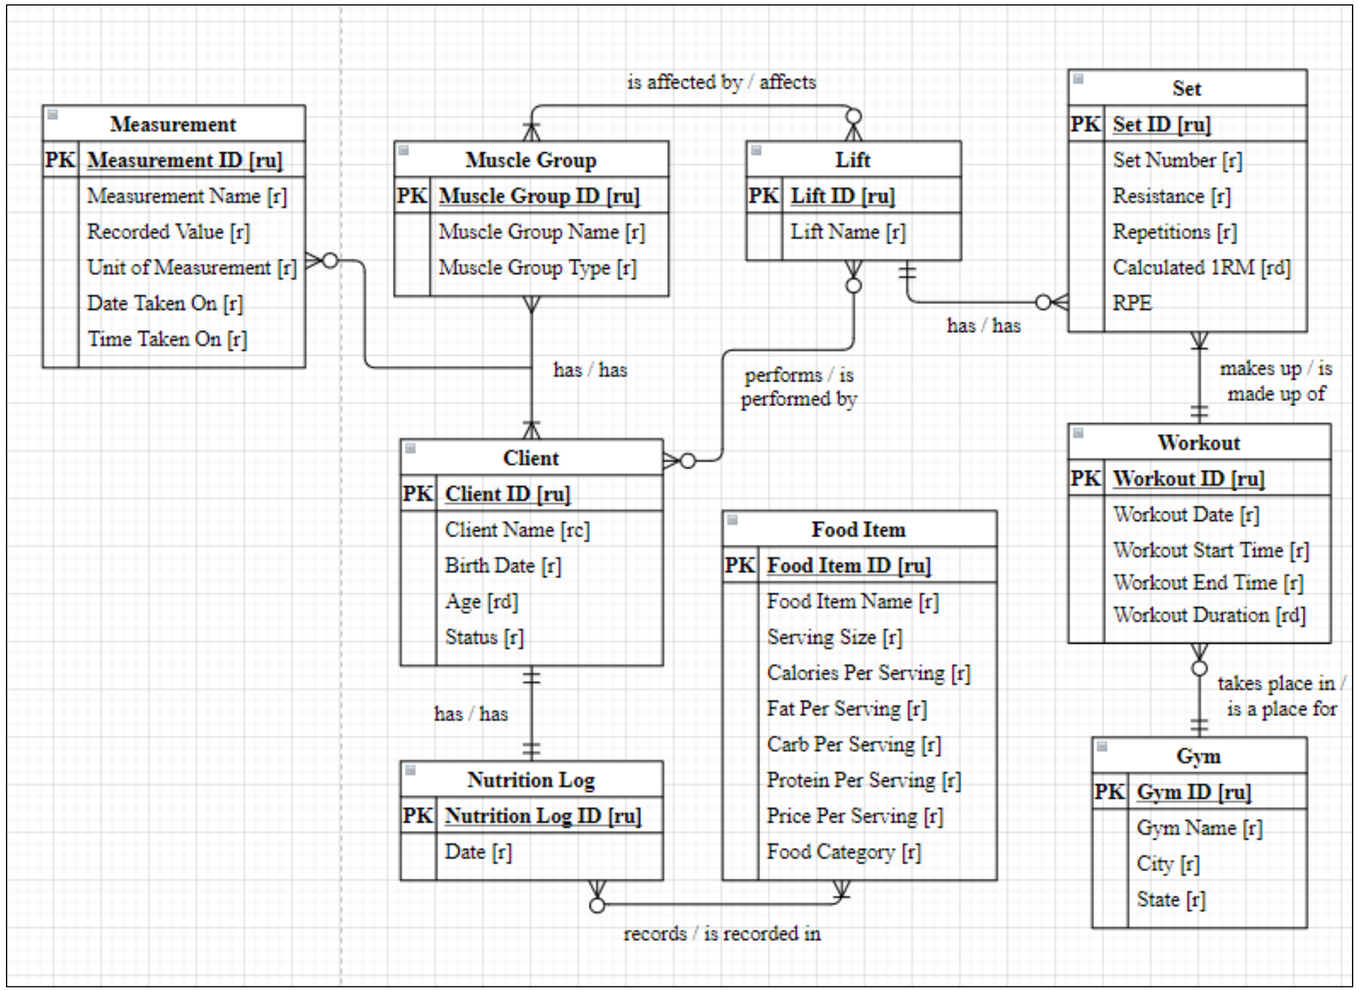

Progress Dashboard

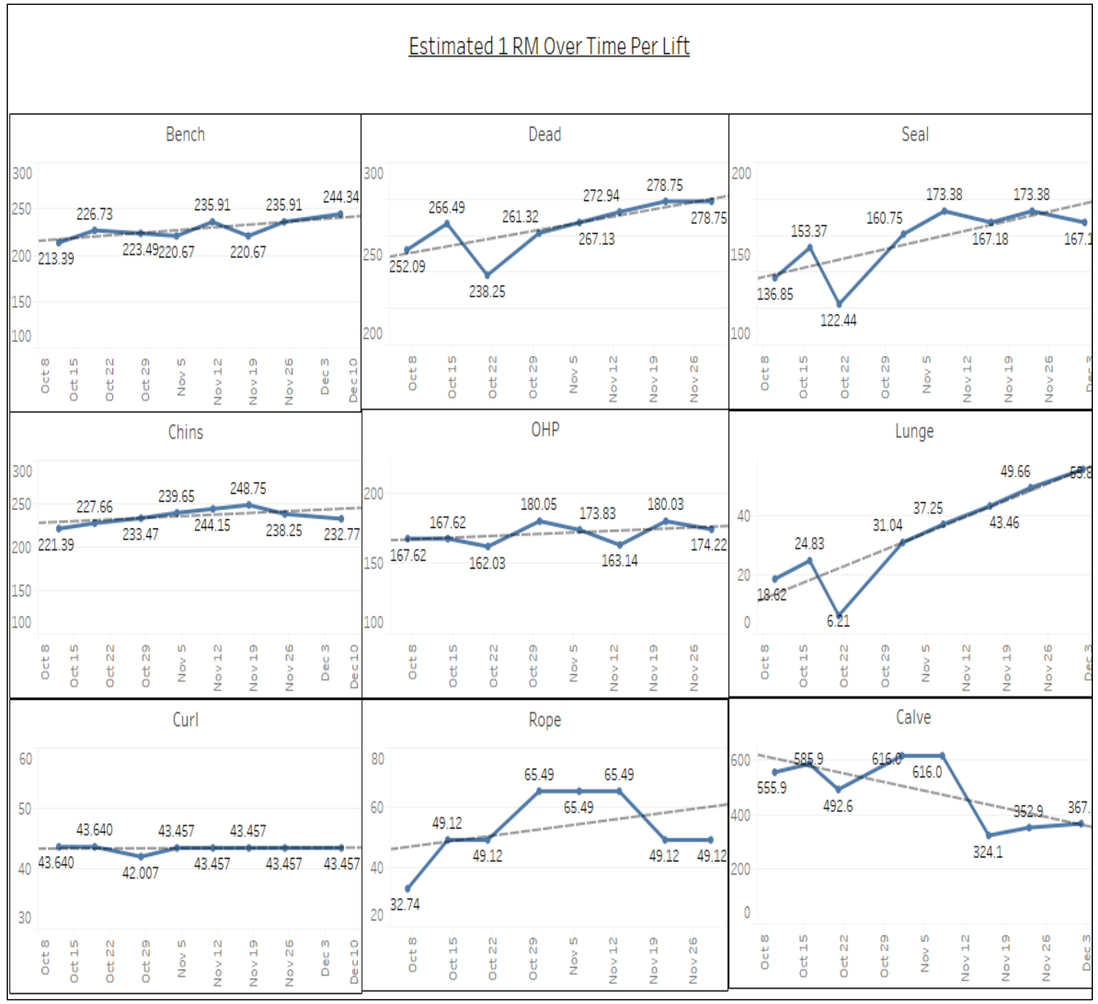

Starting Weight	Current Weight
as of: 10/5/2020	as of: 12/7/2020
167.9	177.0

Week over Week Change

Date Taken On	Recorded Value	% Difference
10/5/2020	167.9	
10/12/2020	166.4	-0.9%
10/19/2020	170.0	2.2%
10/26/2020	169.0	-0.6%
11/2/2020	167.7	-0.8%
11/9/2020	172.1	2.6%
11/16/2020	178.5	3.7%
11/23/2020	173.2	-3.0%
11/30/2020	180.7	4.3%
12/7/2020	177.0	-2.0%

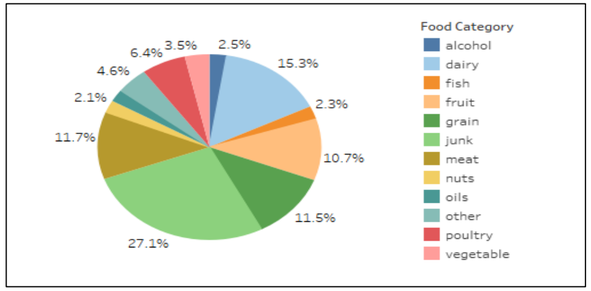

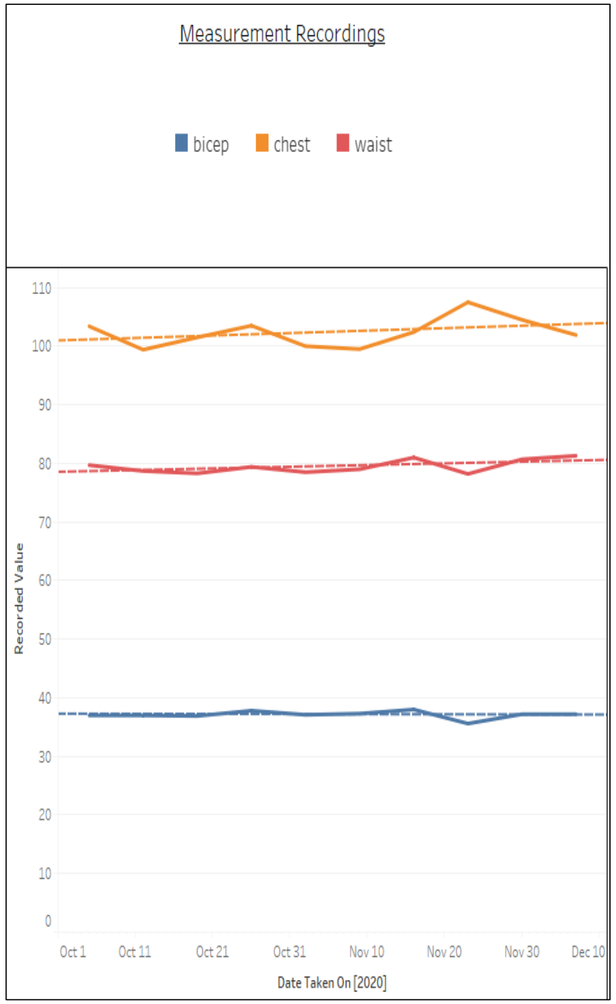

III. Project 2

Machine Learning with Song Data

Part 1: Predicting Song Popularity

Goal: predict the popularity of songs based on audio features and other meta data

Real world application

- Data science is used in developing songs
- How do record labels produce hit song after hit song?
- They have figured out the “formula”

Questions to answer

- How have popular songs change overtime?
- Are there certain attributes that correlate with popular songs?
- Can the popularity of a song be predicted based on its attributes?

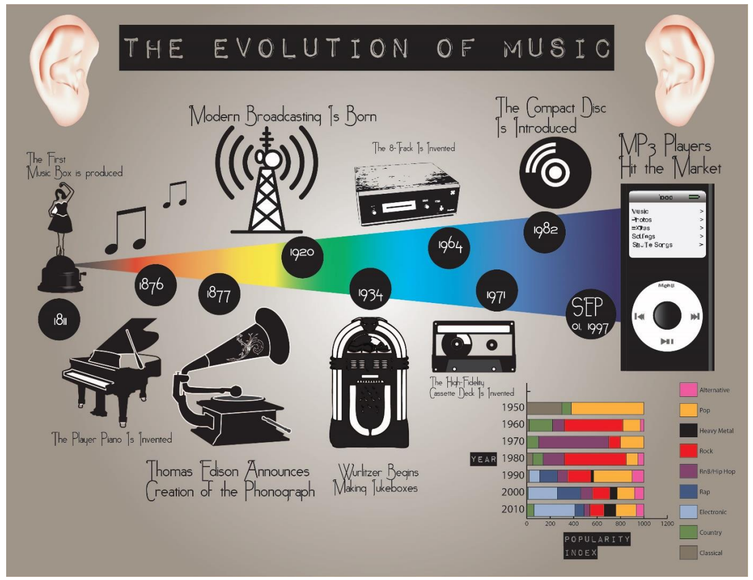

Data Collection

Original dataset from Kaggle and then collected additional metadata through web scraping Wikipedia

Database diagram

Data source 1

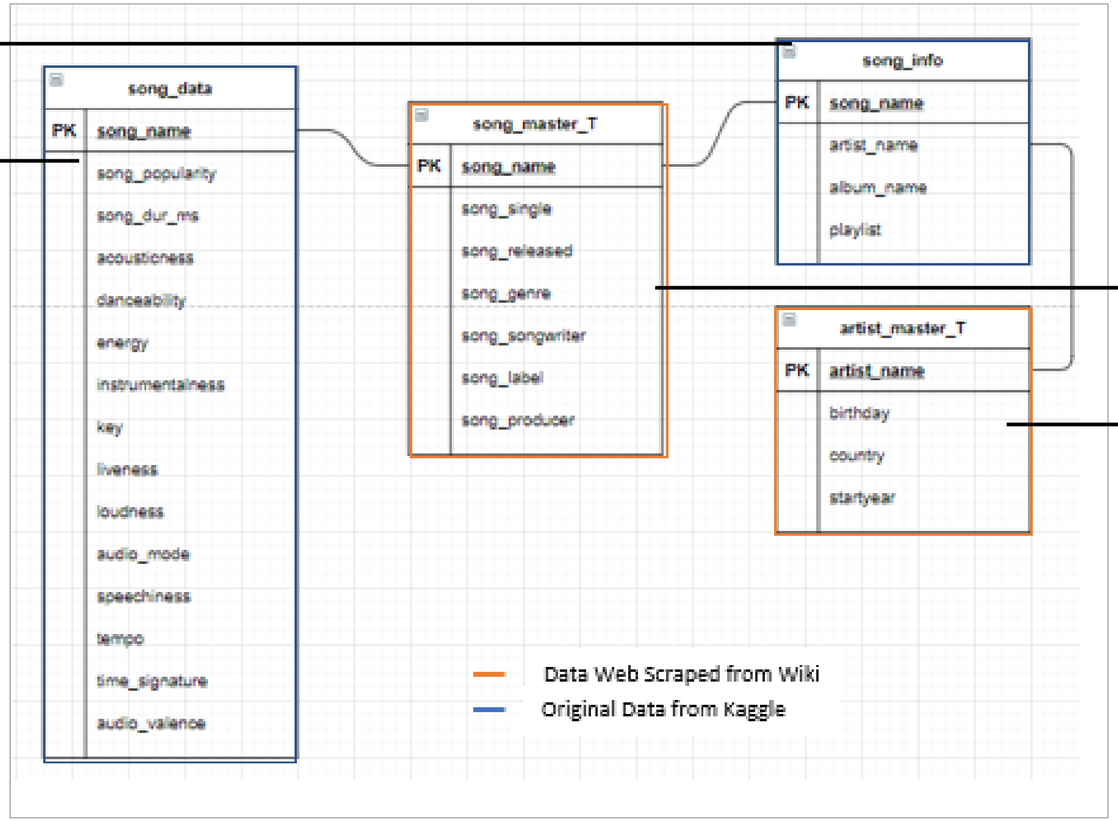

Data source 2

Exploratory Data Analysis

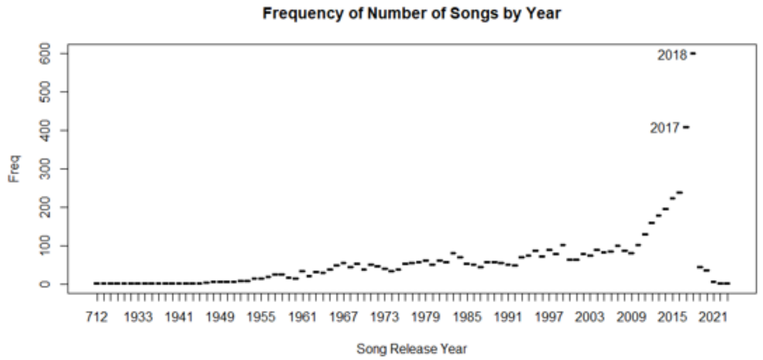

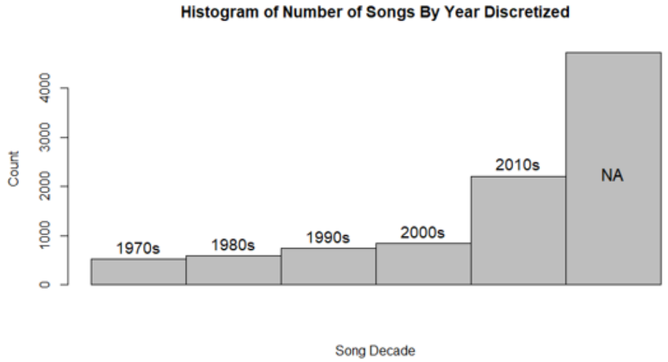

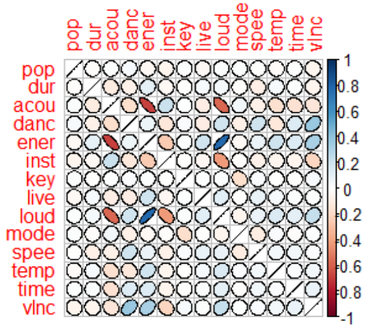

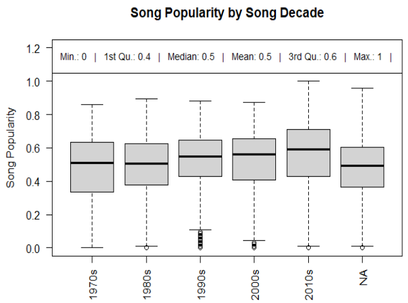

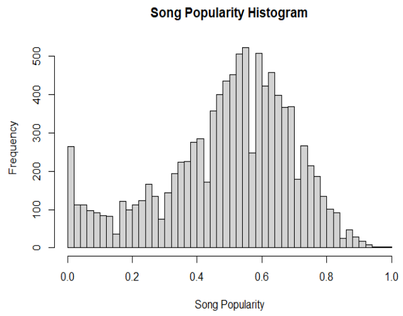

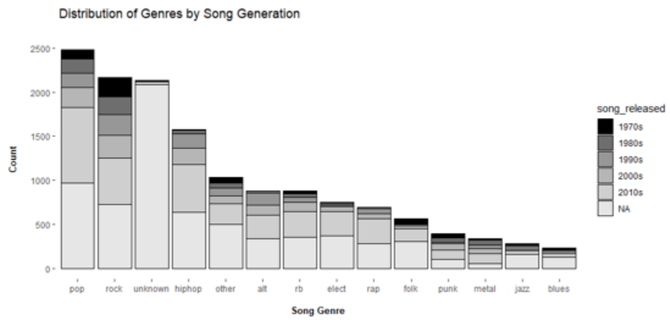

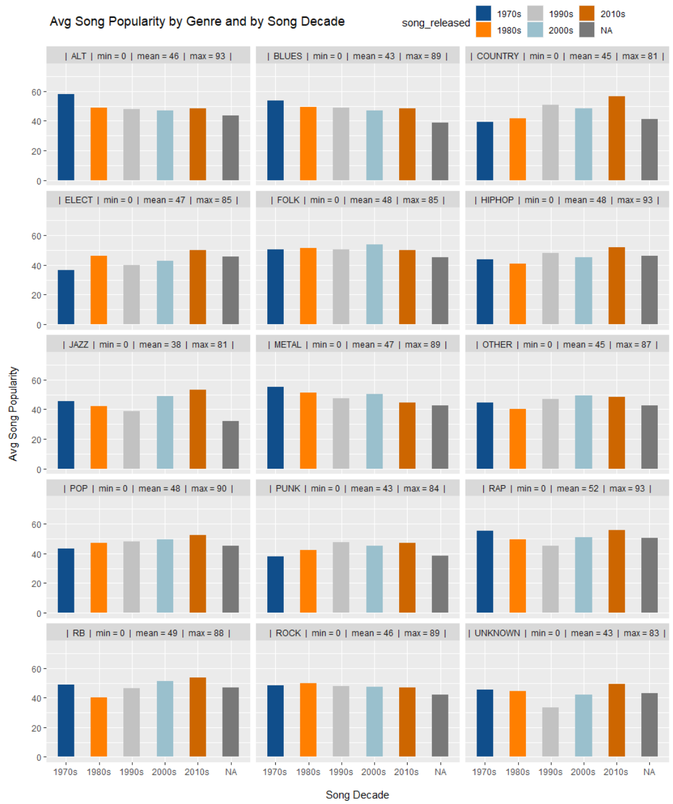

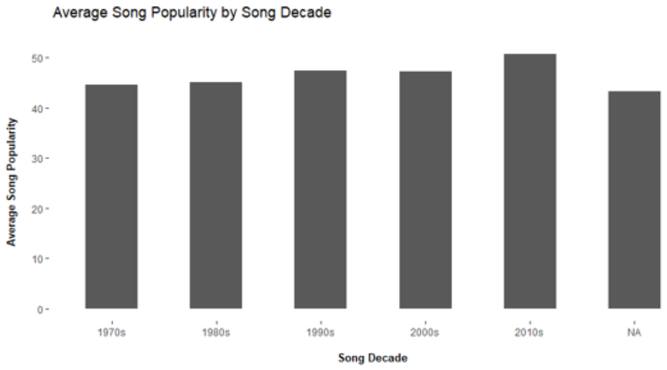

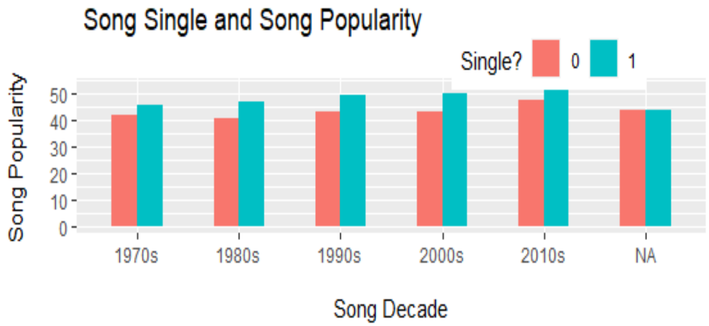

Data Modeling

Type of problem: supervised regression

Song Audio features

Other Song Metadata

Random forest model

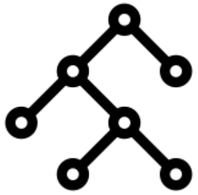

Results

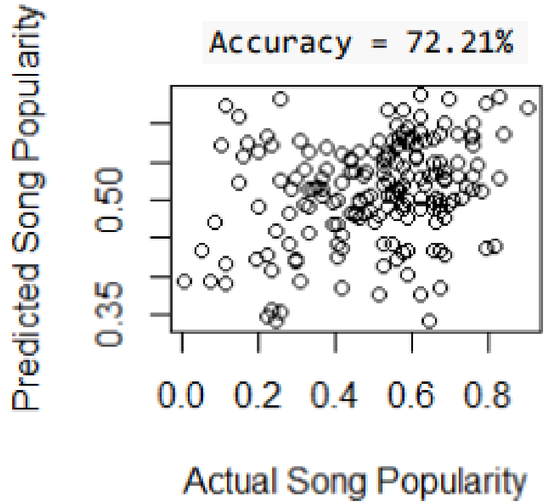

Part 2: Creating Song Clusters

Goal: create clusters of songs based on their attributes

Real world application

- Music streaming services use machine learning to make recommendations
- Auto generating playlists
- Discovering new music

Questions to answer

- Is it possible to create song clusters based on their attributes?
- How do the clusters compare to genres?

Data Collection

Sample random song from Spotify API
- Over 12,000 “pseudo” random songs via Spotify API
- There is no true random method for doing this
- Songs dating back from 1970 to present

Get audio features from Spotify API
- Same audio features from part 1

Get the link to the song on Genius
- Find the correct link by web scraping a google search
- Conduct search for [Song Name] + “Genius.com”
- If there is one it appears in the top results

Get the lyrics for the song from Genius
- Web scrape the HTML code for the genius page
- Parse it to extract the text of the song lyrics

Exploratory Data Analysis

Histogram of song years

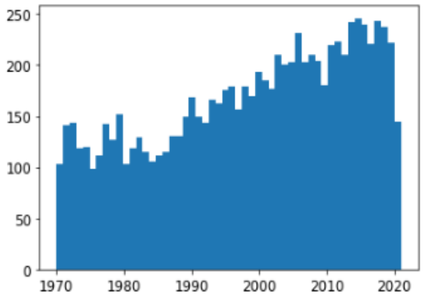

Histogram of song danceability

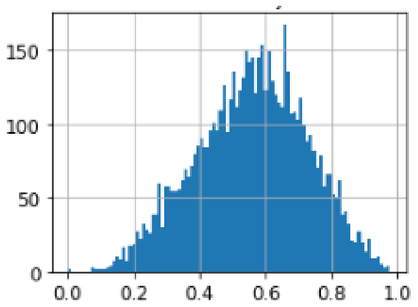

Histogram of song length

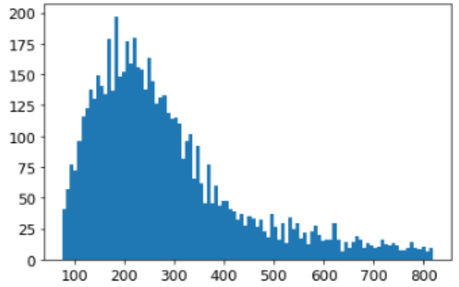

Histogram of song tempo

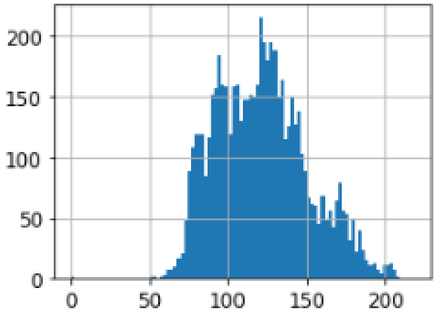

Histogram of polarity scores

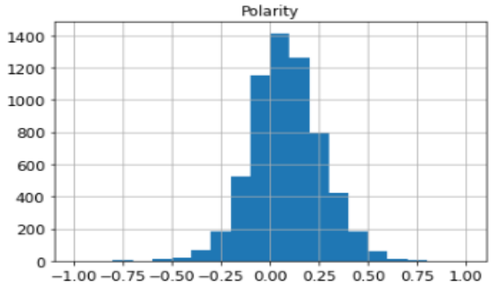

Histogram of song loudness

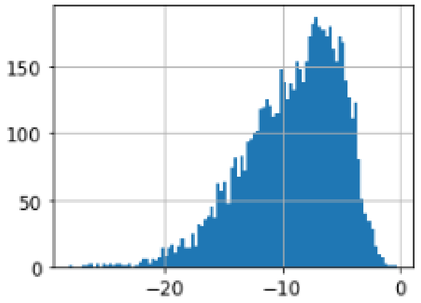

LDA Topic Modeling Results

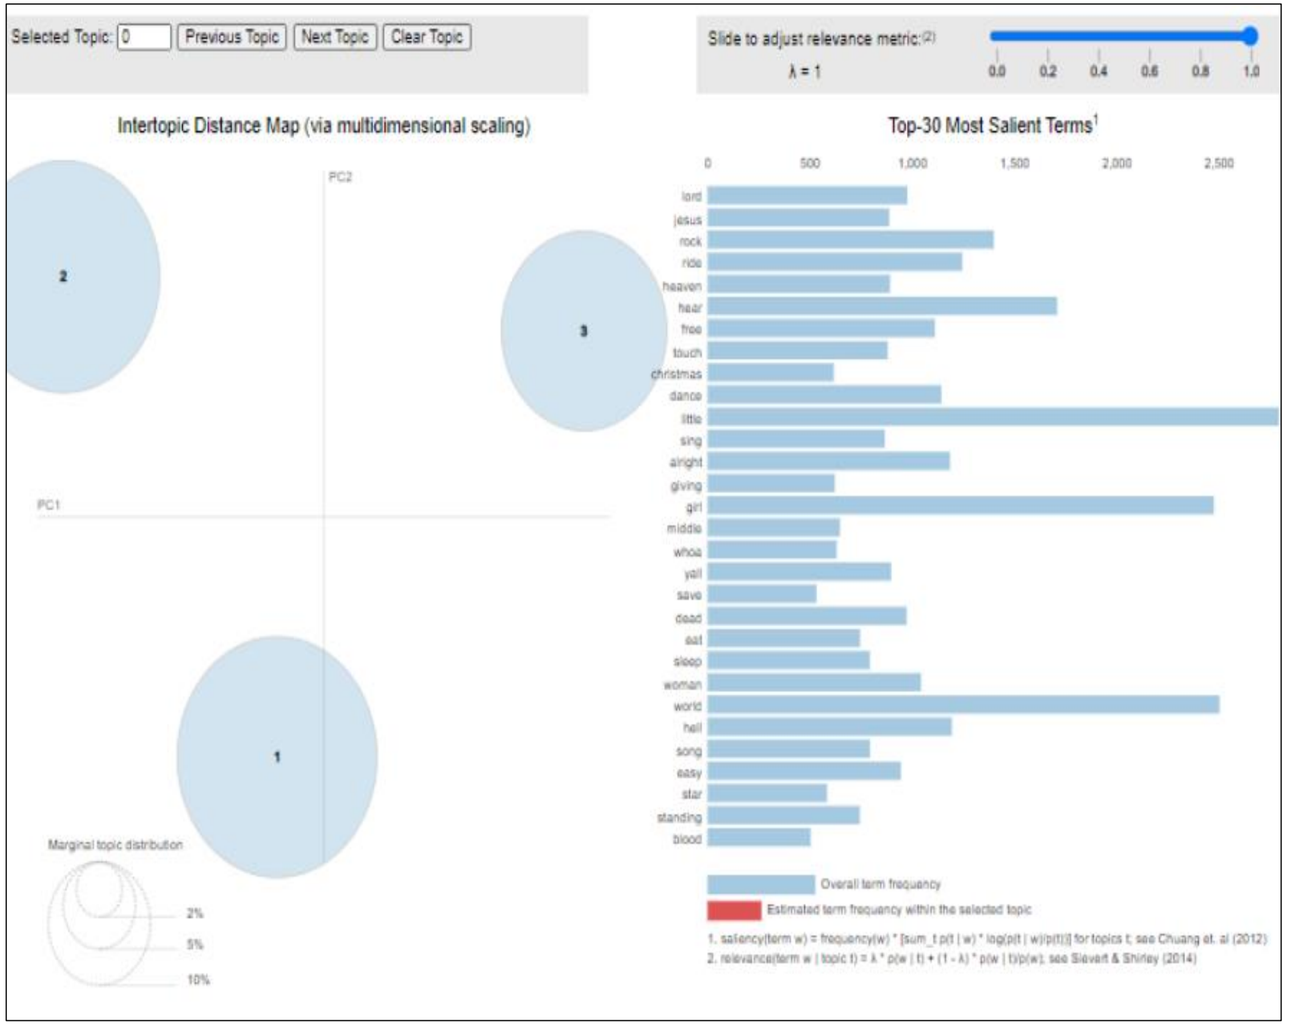

Data Modeling

Type of problem: unsupervised clustering

How the song sounds (audio features)

What the song is about (LDA topic modeling)

Are the song lyrics positive or negative?

K means clustering algorithm

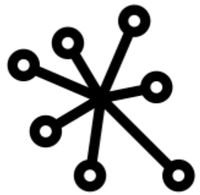

Song clusters k = 7

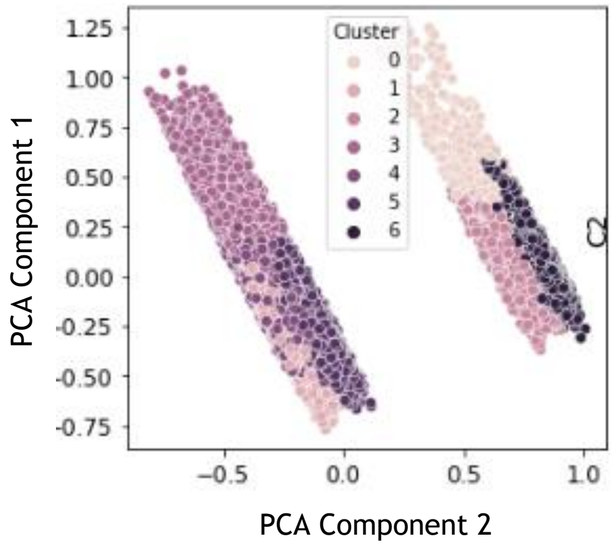

IV. Project 3

Cryptocurrency Price Prediction

Data Collection

Asset table
- Collected by web scraping Coinbase asset directory page
- This page has a list of all tradeable cryptocurrencies on Coinbase
- Currently about 165 cryptocurrencies that can be traded

Price table
- Collected through the yfinance Python package
- Historical price and volume data by day for all data available on Yahoo Finance

CoinMarket table
- Collected through the CoinMarketCap API
- Daily snapshots of circulating supply, max supply, and coin market cap rank
- Cannot get historical data can only take the daily snapshots

Youtube table
- Collected through Google Cloud / Youtube API
- Retrieve a count of published Youtube videos for crypto slugs by day
- Capped at 1000 per day (100 for 10 API keys), but can get historical data
- Starting from the top down and building the history up overtime

Twitter table
- Collected through Twitter API
- Retrieve a count of Tweets that contain hashtag of the ticker by day
- Cannot get historical data without elevated access. Snapshot taken every day

Stocks table
- Collected through yfinance Python package
- Historical closing prices for all data available on Yahoo Finance

Google table
- Collected through pytrends python package
- Retrieve data from Google Trends about cryptocurrency searches by day
- Might need to use some extrapolation methods due to data availability

Database diagram

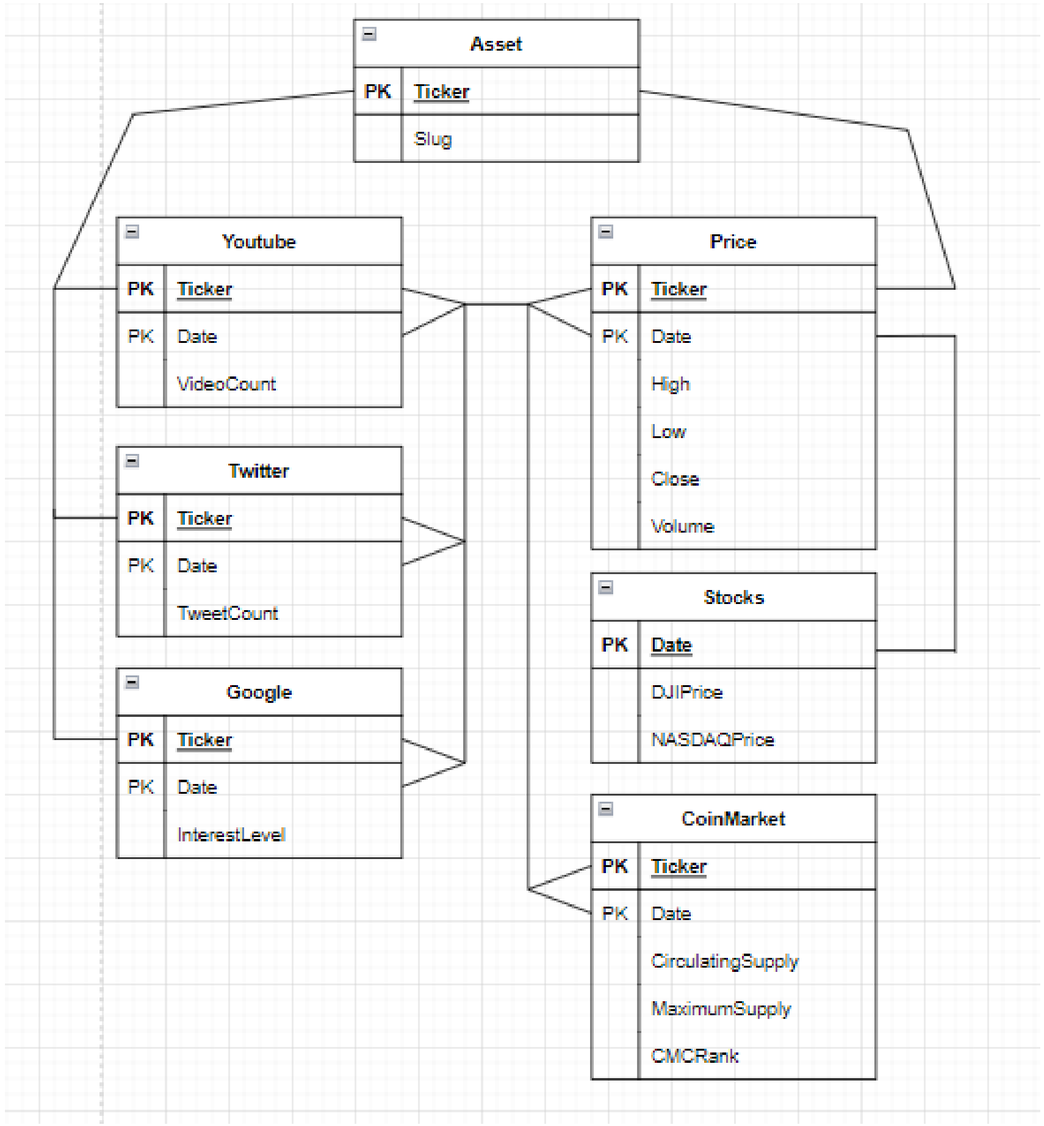

Exploratory Data Analysis

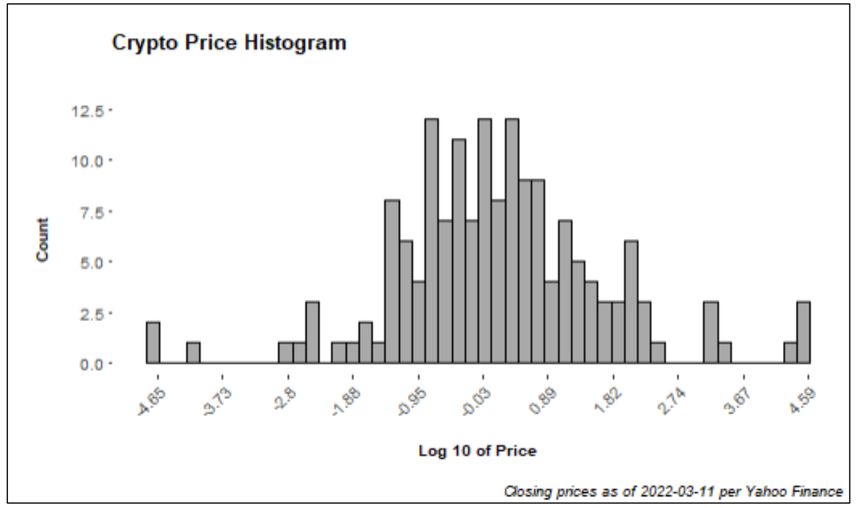

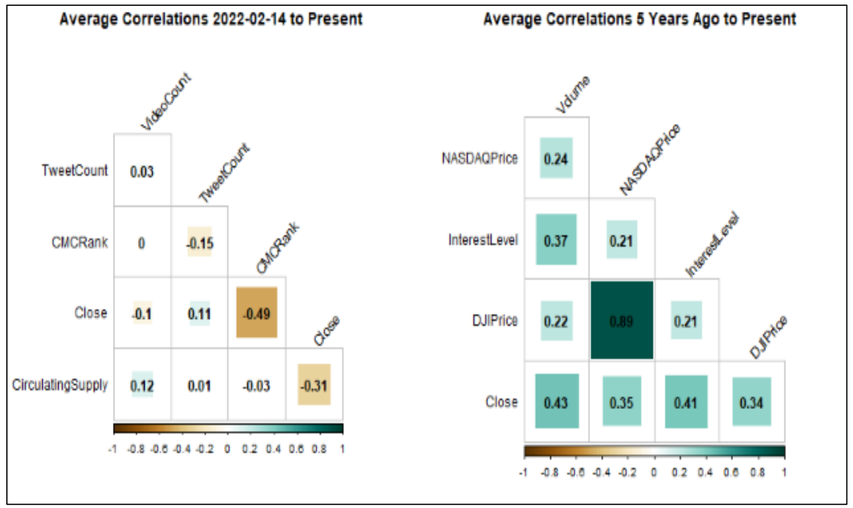

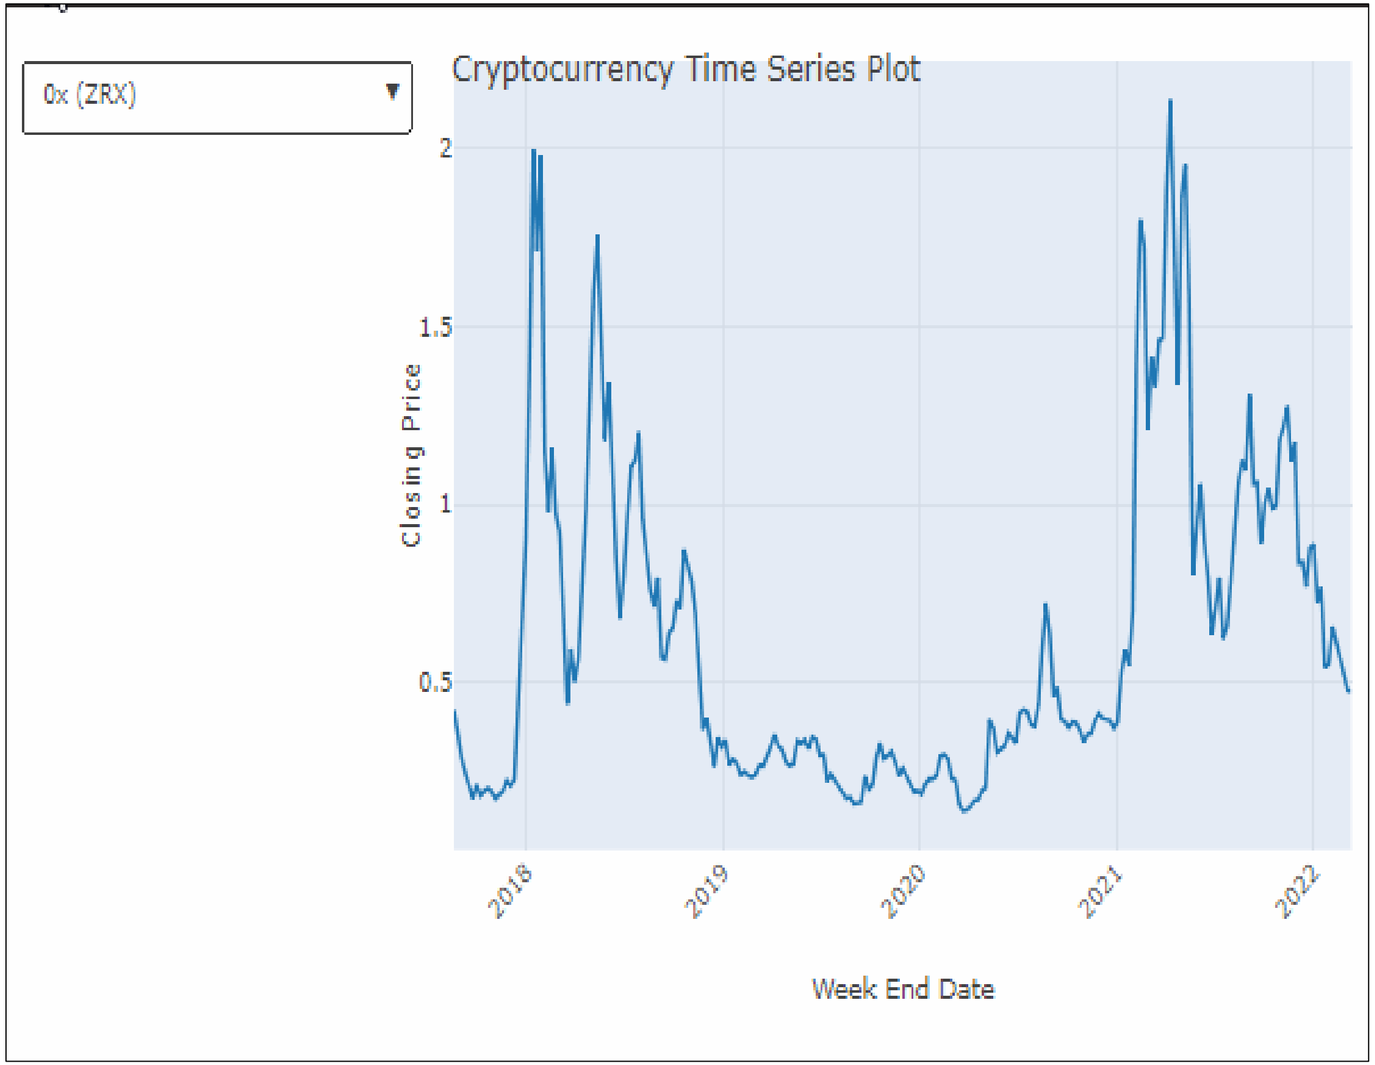

V. Conclusion

Reflection

Project I

Flexibility as a point of emphasis
- Once the database is developed it can be difficult to make changes
- But changes are bound to happen, so flexibility is important
- Discuss with stakeholders, put yourself in their shoes

Project II Part 1

Importance of data quality
- There is only so much that can be done to remedy bad data
- If the data that goes into a model is bad the data that comes out is also likely bad
- If the data scientist is involved in data collection, make sure its done right!

Project II Part 2

Dealing with ambiguity
- Ambiguity is a reoccurring theme in data science, especially with text data
- Crowdsourcing is a good way to approach ambiguity
- Do not expect out of the box solutions to be enough, configure for the task at hand

Gallery

A collection of other data visualizations I have created for assignments, labs, or projects outside of the 3 projects discussed in this presentation

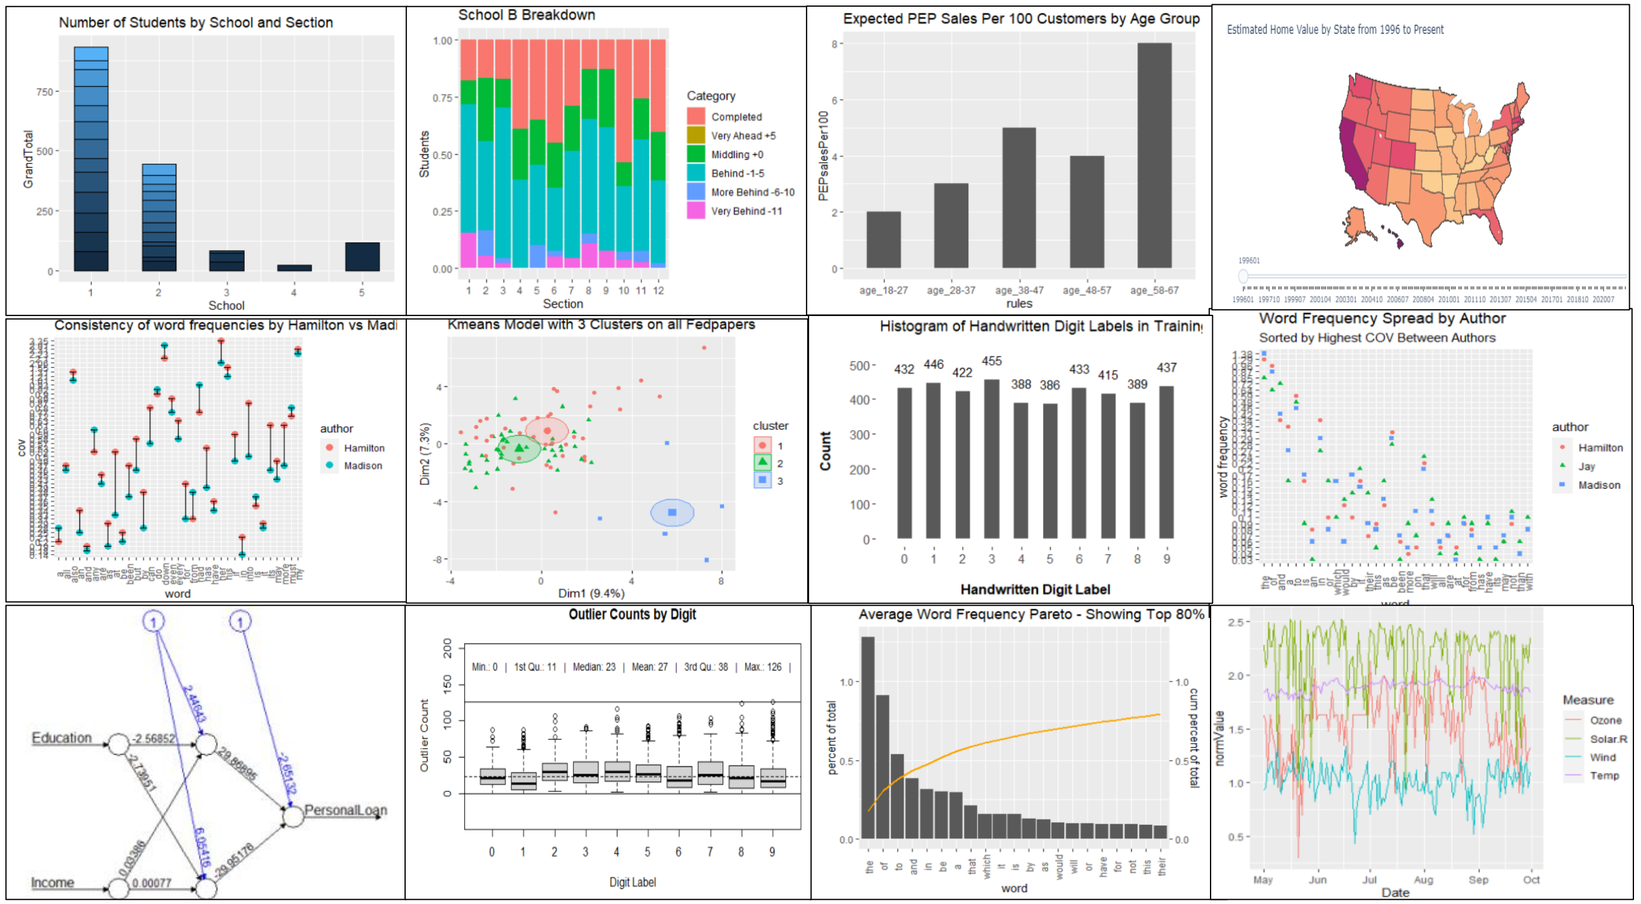

End Presentation